

GVPT Data Science Camp

Data Visualization

Ryan Bookstein

August 25, 2026

University of Maryland

- ▶ Introduction to using R and RStudio

Learning Objectives

- ▶ Introduction to using `R` and `RStudio`
- ▶ Create your first plot in `R`

Learning Objectives

- ▶ Introduction to using **R** and **RStudio**
- ▶ Create your first plot in **R**
- ▶ Test your hypotheses using data visualizations

- ▶ Data visualization creates informative displays of data

- ▶ Data visualization creates informative displays of data
- ▶ We can display data in ways beyond what we will learn today

- ▶ Data visualization creates informative displays of data
- ▶ We can display data in ways beyond what we will learn today
- ▶ Data visualization is an important component of a research project

R Basics

- ▶ R evaluates the code you type — try some arithmetic:

```
1 + 2
```

- ▶ R evaluates the code you type — try some arithmetic:

```
1 + 2
```

- ▶ Functions take arguments inside parentheses:

```
sum(1, 2)
```

Exercise

- ▶ Find the sum of 412, 65, and 292

Exercise

- ▶ Find the sum of 412, 65, and 292
- ▶ What is 48 divided by 5?

- ▶ Find the sum of 412, 65, and 292
- ▶ What is 48 divided by 5?
- ▶ What is the square root of 5232? HINT: use the `sqrt()` function

- Two ways to get the sum:

```
sum(412, 65, 292)
```

```
#> [1] 769
```

```
412 + 65 + 292
```

```
#> [1] 769
```

- ▶ Two ways to get the sum:

```
sum(412, 65, 292)  
#> [1] 769  
412 + 65 + 292  
#> [1] 769
```

- ▶ Division:

```
48 / 5  
#> [1] 9.6
```

Check Your Answers

- ▶ Two ways to get the sum:

```
sum(412, 65, 292)
#> [1] 769
412 + 65 + 292
#> [1] 769
```

- ▶ Division:

```
48 / 5
#> [1] 9.6
```

- ▶ Square root:

```
sqrt(5232)
#> [1] 72.33257
```


- ▶ R has a large number of built-in functions that allow you to manipulate data

- ▶ R has a large number of built-in functions that allow you to manipulate data
- ▶ For example, the `seq()` function provides all numbers in a specified sequence:

```
seq(from = 1, to = 20, by = 3)  
#> [1] 1 4 7 10 13 16 19
```

- ▶ R has a large number of built-in functions that allow you to manipulate data
- ▶ For example, the `seq()` function provides all numbers in a specified sequence:

```
seq(from = 1, to = 20, by = 3)  
#> [1]  1  4  7 10 13 16 19
```

- ▶ Arguments are positional, so naming them is optional:

```
seq(1, 20, 3)  
#> [1]  1  4  7 10 13 16 19
```

- ▶ Packages are collections of R functions and data

- ▶ Packages are collections of R functions and data
- ▶ Install the relevant package(s) once:

```
install.packages("tidyverse")
```

- ▶ Packages are collections of R functions and data
- ▶ Install the relevant package(s) once:

```
install.packages("tidyverse")
```

- ▶ Load the packages in your current session:

```
library(tidyverse)
```

- ▶ Using the RStudio console, install the `tidyverse` packages

```
install.packages("tidyverse")
```

- ▶ Using the RStudio console, install the `tidyverse` packages

```
install.packages("tidyverse")
```

- ▶ Load these packages in your current session, along with `readxl` (installed with the tidyverse) for reading Excel files

```
library(tidyverse)  
library(readxl)
```


- ▶ *For you and your co-author's sake*

- ▶ *For you and your co-author's sake*
- ▶ Keep everything:


- ▶ *For you and your co-author's sake*
- ▶ Keep everything:
 - ▷ Organized

- ▶ *For you and your co-author's sake*
- ▶ Keep everything:
 - ▷ Organized
 - ▷ Reproducible

- ▶ *For you and your co-author's sake*
- ▶ Keep everything:
 - ▷ Organized
 - ▷ Reproducible
 - ▷ Annotated


In the RStudio console, type `getwd()` to locate where you are in your computer

Create a new RStudio project for this Data Science Camp




Source: *R for Data Science* (R4DS)

Create a new RStudio project for this Data Science Camp



Source: *R for Data Science* (R4DS)

Create a new RStudio project for this Data Science Camp




Source: *R for Data Science* (R4DS)

Re-run `getwd()` to locate where you are now in your computer

Introduction to Data Visualization

- ▶ Data visualization is the graphical representation of information and data

- ▶ Data visualization is the graphical representation of information and data
- ▶ From the Center for Effective Lawmaking (CEL) data: Do more senior members of Congress pass more legislation than newer members?



How Did We Do This?

R code:

```
ggplot(data = cel, mapping = aes(x = seniority, y = les)) +  
  geom_point(mapping = aes(color = party), alpha = 0.15,  
             position = position_jitter(width = 0.25)) +  
  scale_color_manual(name = "Party", values = party_colors) +  
  geom_smooth(method = "lm") +  
  theme(legend.position = "bottom",  
        panel.grid = element_blank(),  
        panel.background = element_blank(),  
        plot.title.position = "plot",  
        plot.title = element_text(face = "bold")) +  
  labs(title = "Seniority and Legislative Effectiveness",  
        subtitle = "U.S. House Members, 1973-2023, by Party",  
        x = "Seniority (Terms Served)",  
        y = "Legislative Effectiveness Score (LES)")
```

Setting Up the CEL Data

- ▶ Load the relevant packages:

```
library(tidyverse)
library(readxl)
```

Load Relevant Packages and Data

- ▶ Load the relevant packages:

```
library(tidyverse)
library(readxl)
```

- ▶ Download the Center for Effective Lawmaking House dataset ([CELHouse93to118](#)) from thelawmakers.org

Load Relevant Packages and Data

- ▶ Load the relevant packages:

```
library(tidyverse)
library(readxl)
```

- ▶ Download the Center for Effective Lawmaking House dataset ([CELHouse93to118](#)) from thelawmakers.org
- ▶ Save it in your project folder as [CELHouse93to118.xlsx](#)

- ▶ Read the Excel file into an object named `cel`:

```
cel <- read_excel("CELHouse93to118.xlsx")
```

- ▶ Read the Excel file into an object named `cel`:

```
cel <- read_excel("CELHouse93to118.xlsx")
```

- ▶ One row per member per Congress: every member of the U.S. House from the 93rd Congress (1973) through the 118th (2023)

Rename the Variables

The columns arrive with long descriptive names — give them short ones. Copy this chunk as is; we will unpack how it works tomorrow:

```
cel <- rename(cel,  
  member      = "Legislator name, as given in THOMAS",  
  congress    = "Congress number",  
  year        = "Year at start of Congress",  
  state       = "Two-letter state code",  
  dem         = "1 = Democrat",  
  female      = "1 = female",  
  vote_pct    = "Percent vote received to enter this Congress",  
  ideology    = "First-dimension DW-NOMINATE score",  
  seniority   = "Seniority, number of terms served counting current",  
  born        = "Year born",  
  les         = "LES 1.0",  
  bills_passed = "Total number of standalone bills passing House",  
  benchmark   = "1 = Below, 2 = Meets, 3 = Above, Based on Classic LES")
```

Keep only the renamed columns, and add readable labels for party and for the CEL benchmark categories:

```
cel <- select(cel, member, congress, year, state, dem, female,
              vote_pct, ideology, seniority, born, les,
              bills_passed, benchmark)

cel <- mutate(cel,
  party = if_else(dem == 1, "Democrat", "Republican"),
  benchmark = factor(benchmark,
                     labels = c("Below", "Meets", "Above")))
```

- ▶ Learn more about this dataset by browsing the CEL codebook at thelawmakers.org

- ▶ Learn more about this dataset by browsing the CEL codebook at thelawmakers.org
- ▶ Get a quick overview of every variable in your console:

```
glimpse(CEL)
```

- ▶ Open the data in a spreadsheet-style viewer:

```
view(cel)
```

- ▶ Open the data in a spreadsheet-style viewer:

```
view(ce1)
```

- ▶ A couple of useful variables — **seniority**: number of terms served, counting the current one; **les**: Legislative Effectiveness Score, a summary of how successful a member is at moving their bills forward

- How many rows and columns are in `cel`?

```
nrow(cel)
```

```
ncol(cel)
```

- ▶ How many rows and columns are in `cel`?

```
nrow(cel)
ncol(cel)
```

- ▶ What does the `vote_pct` variable describe?

Building the Plot


Set Up Your Plot

```
ggplot(data = cel)
```




Map Your Aesthetics

```
ggplot(data = cel, mapping = aes(x = seniority, y = les))
```



Add in Your Legislators

```
ggplot(data = cel, mapping = aes(x = seniority, y = les)) +  
  geom_point()
```



- ▶ Inspect the `les` and `vote_pct` variables and describe them

- ▶ Inspect the `les` and `vote_pct` variables and describe them
- ▶ Make a scatter plot of them

- ▶ Why does the following code give an error, and how would you fix it?

```
ggplot(data = cel) +  
  geom_point()
```

Let's Look at Groups in the Data


- ▶ You can use visual elements or aesthetics (`aes`) to communicate your data

Let's Look at Groups in the Data

- ▶ You can use visual elements or aesthetics (`aes`) to communicate your data
- ▶ Let's look at a categorical variable: the party of the member (Democrat or Republican)

Let's Plot Groups in the Data

```
ggplot(data = cel,  
       mapping = aes(x = seniority, y = les, color = party)) +  
  geom_point(alpha = 0.15,  
            position = position_jitter(width = 0.25))
```



Two Problems With That Plot

- ▶ `cel` has 11,606 rows — one per member per Congress — so the points sit on top of each other. `alpha` and `position_jitter()` let us see through the pile


Two Problems With That Plot

- ▶ `cel` has 11,606 rows — one per member per Congress — so the points sit on top of each other. `alpha` and `position_jitter()` let us see through the pile
- ▶ R chose the colours alphabetically, so Democrats came out red and Republicans blue — we should choose them ourselves

Choosing Your Own Colours


```
party_colors <- c(Democrat = "#1B5FA8", Republican = "#C0392B")

ggplot(data = cel,
       mapping = aes(x = seniority, y = les, color = party)) +
  geom_point(alpha = 0.15,
             position = position_jitter(width = 0.25)) +
  scale_color_manual(values = party_colors)
```




Add More Information

```
ggplot(data = cel,  
       mapping = aes(x = seniority, y = les, color = party)) +  
  geom_point(alpha = 0.15,  
            position = position_jitter(width = 0.25)) +  
  scale_color_manual(values = party_colors) +  
  geom_smooth(method = "lm", se = FALSE)
```



Look at the Relationship Among All Members


```
ggplot(data = cel, mapping = aes(x = seniority, y = les)) +  
  geom_point(mapping = aes(color = party), alpha = 0.15,  
            position = position_jitter(width = 0.25)) +  
  scale_color_manual(values = party_colors) +  
  geom_smooth(method = "lm", se = FALSE)
```



Polishing the Plot


Add Titles and Labels

```
ggplot(data = cel, mapping = aes(x = seniority, y = les)) +  
  geom_point(mapping = aes(color = party), alpha = 0.15,  
             position = position_jitter(width = 0.25)) +  
  scale_color_manual(name = "Party", values = party_colors) +  
  geom_smooth(method = "lm", se = FALSE) +  
  labs(  
    title = "Seniority and Legislative Effectiveness",  
    subtitle = "U.S. House Members, 1973-2023, by Party",  
    x = "Seniority (Terms Served)",  
    y = "Legislative Effectiveness Score (LES)"  
  )
```




```
ggplot(cel, aes(x = seniority, y = les)) +  
  geom_point(data = filter(cel, state != "MD"),  
            color = "grey85", alpha = 0.4) +  
  geom_point(data = filter(cel, state == "MD"),  
            color = "#C0392B")
```

- Only 211 of the 11,606 rows are Maryland members, so we draw everyone else in grey first and put Maryland on top



Editing Plots

```
ggplot(data = cel) +  
  geom_point(mapping = aes(x = seniority, y = les), color = "red")
```



- ▶ What is wrong with this code?


```
ggplot(data = cel) +  
  geom_point(mapping = aes(x = seniority, y = les, color = "blue"))
```

- ▶ Name a categorical and continuous variable in `cel`

- ▶ Name a categorical and continuous variable in `cel`
- ▶ Map a continuous variable (try `votepct`) to color. How do the aesthetics behave differently for categorical and continuous variables?


- ▶ Name a categorical and continuous variable in `cel`
- ▶ Map a continuous variable (try `votepct`) to color. How do the aesthetics behave differently for categorical and continuous variables?
- ▶ Map `state` to the shape aesthetic. What does the warning say?

```
ggplot(data = cel, mapping = aes(x = seniority, y = les)) +  
  geom_point(mapping = aes(color = party), alpha = 0.15,  
             position = position_jitter(width = 0.25)) +  
  scale_color_manual(name = "Party", values = party_colors) +  
  geom_smooth(method = "lm", se = FALSE) +  
  theme_minimal() +  
  labs(  
    title = "Seniority and Legislative Effectiveness",  
    subtitle = "U.S. House Members, 1973-2023, by Party",  
    x = "Seniority (Terms Served)",  
    y = "Legislative Effectiveness Score (LES)"  
  )
```



Creating a Unique Theme

```
ggplot(data = cel, mapping = aes(x = seniority, y = les)) +  
  geom_point(mapping = aes(color = party), alpha = 0.15,  
             position = position_jitter(width = 0.25)) +  
  scale_color_manual(name = "Party", values = party_colors) +  
  geom_smooth(method = "lm") +  
  theme(legend.position = "bottom",  
        panel.grid = element_blank(),  
        panel.background = element_blank(),  
        plot.title.position = "plot",  
        plot.title = element_text(face = "bold")) +  
  labs(title = "Seniority and Legislative Effectiveness",  
        subtitle = "U.S. House Members, 1973-2023, by Party",  
        x = "Seniority (Terms Served)",  
        y = "Legislative Effectiveness Score (LES)")
```


- ▶ Customize the scatter plot you created using the `theme()` argument

Other Chart Types


- ▶ When visualizing and reporting categorical data, we must use different tools

```
select(cel, member, state, party)
```

```
ggplot(cel, aes(x = benchmark)) +  
  geom_bar() +  
  labs(x = "CEL Benchmark (Classic LES)", y = "Number of Members")
```




```
ggplot(cel, aes(x = fct_infreq(benchmark))) +  
  geom_bar() +  
  labs(x = "CEL Benchmark (Classic LES)", y = "Number of Members")
```




Visualizing Numerical Data

```
ggplot(cel, aes(x = vote_pct)) +  
  geom_histogram(binwidth = 2) +  
  labs(x = "Percent of Vote Received", y = "Number of Members")
```




Visualizing Numerical Data

```
ggplot(cel, aes(x = vote_pct)) +  
  geom_density() +  
  labs(x = "Percent of Vote Received", y = "Density")
```




Visualizing Numerical Data

```
ggplot(cel, aes(x = ideology, color = party)) +  
  geom_density() +  
  scale_color_manual(name = "Party", values = party_colors) +  
  labs(x = "DW-NOMINATE (First Dimension)", y = "Density")
```




Visualizing Numerical Data

```
ggplot(cel, aes(x = ideology, color = party, fill = party)) +  
  geom_density(alpha = 0.5) +  
  scale_color_manual(name = "Party", values = party_colors) +  
  scale_fill_manual(name = "Party", values = party_colors) +  
  labs(x = "DW-NOMINATE (First Dimension)", y = "Density")
```



Comparing a Continuous Variable Across Groups

```
ggplot(cel, aes(x = les, y = benchmark)) +  
  geom_boxplot() +  
  labs(x = "Legislative Effectiveness Score (LES)",  
       y = "CEL Benchmark (Classic LES)")
```



- ▶ Set up data science tools

- ▶ Set up data science tools
- ▶ Plotted data in multiple ways

- ▶ Set up data science tools
- ▶ Plotted data in multiple ways
- ▶ Discovered relationships in the data

- ▶ Set up data science tools
- ▶ Plotted data in multiple ways
- ▶ Discovered relationships in the data
- ▶ Connected these relationships to expectations (hypotheses)

- ▶ On Thursday, you will apply the skills learned in this course to an idea that interests you. In doing so, you will need to find a dataset that:

- ▶ On Thursday, you will apply the skills learned in this course to an idea that interests you. In doing so, you will need to find a dataset that:
 - ▶ Is relevant to your research interests

- ▶ On Thursday, you will apply the skills learned in this course to an idea that interests you. In doing so, you will need to find a dataset that:
 - ▶ Is relevant to your research interests
 - ▶ Contains continuous and categorical variables

- ▶ On Thursday, you will apply the skills learned in this course to an idea that interests you. In doing so, you will need to find a dataset that:
 - ▷ Is relevant to your research interests
 - ▷ Contains continuous and categorical variables
- ▶ If needed, reach out to me for dataset recommendations

- ▶ Zinovyev, Andrei. 2010. "Data Visualization in Political and Social Sciences." In *International Encyclopedia of Political Science*, ed. Bertrand Badie, Dirk Berg-Schlosser, and Leonardo Morlino. Thousand Oaks, CA: SAGE.

Questions?